

Introduction to Ethereum

part 2

Week 1
Lesson 4

Lesson Plan

Review of last lesson

Ethereum storage

Ethereum accounts

Ethereum transactions

EVM languages

Ethereum clients and mining

Connecting to a test network

ETH 2.0

This Week

Monday - Ethereum continued

Tuesday - Remix and the developmental process

Wednesday - Introduction to Solidity

Thursday - Solidity part 2

Friday - Useful libraries for Solidity

Reading List

[DEVCON1: Understanding the Ethereum Blockchain Protocol - Vitalik Buterin](#)

[Mastering Ethereum](#) by Andreas Antonopoulos

[White paper](#)

[Beige Paper](#)


[Yellow Paper](#)

Points from last lesson

- Gas and storage models
- Social aspects of attacks on crypto currencies

Ethereum Data Structures


Ethereum uses Merkle Patricia Tries / Radix Tries for their searching performance and low memory footprint



World and Account State

Account storage contents Trie

A mapping between integer keys (KEC) and integer values (RLP)



Account, $\sigma[\text{address}]$

RLP data structure

nonce, $\sigma[\text{address}]_n$

scalar; the number of transactions sent from this address or, in the case of accounts with associated code, the number of contract-creations made by this account.

balance, $\sigma[\text{address}]_b$

scalar; the number of Wei owned by this address

storageRoot, $\sigma[\text{address}]_s$


256-bit hash of the root node of a Merkle Patricia tree that encodes the storage contents of the account (a mapping between 256-bit integer values), encoded into the trie as a mapping from the Keccak 256-bit hash of the 256-bit integer keys to the RLP-encoded 256-bit integer

codeHash, $\sigma[\text{address}]_c$

Hash of the EVM code of this account - the code that gets executed should this address receive a message call; it is immutable and thus, unlike all other fields, cannot be changed after construction. All such code fragments are contained in the state database under their corresponding hashes for later retrieval.

World State Trie, σ

A mapping between addresses and account states. Stored as a merkle-paritica tree



Transactions

Transaction, T

nonce, T_n

Scalar; the number of transactions sent by the sender

gasPrice, T_p

scalar; the number of Wei to be paid per unit of gas for all computation costs incurred as a result of execution of this transaction

gasLimit, T_g

scalar; the max amount of gas that should be used in executing this transaction. Paid before any computation is done, may not be increased later.

to, T_t


160-bit address of the message call's recipient or, for a contract creation transaction, 0 (zero bytes).

value, T_v

scalar; the number of Wei to be transferred to the message call's recipient or, in the case of contract creation, as an endowment to the newly created account.

Transaction Trie, T

A merkle-paritica tree of transactions to include



Transaction Receipts and Logs

init, T_i Contract creation transaction only; unlimited size byte array specifying the EVM-code for the account initialisation	data, T_d Message call transaction only; unlimited size byte array specifying the input data of the message call
v, r, s, T_w, T_r, T_s Outputs from EDSCA signature of $\text{KEC}(T_o, T_p, T_i, T_r, T_d \text{ or } T_d)$ identifies sender.	

Transaction Receipt, $B_R[i]$
(i =transaction no)
 Tuple

post-transaction state, R_o
 Trie

cumulative gas used, R_u
 positive integer

transaction logs, R_l
 Series of log entries (tuples).

bloom filter of log info, R_b
 256 byte hash. Reduction of a log entry.

Log Entry, O
 Tuple


Logger's address, O_a
 address

Log topics, O_t
 A series of 32-byte log topics (O_{t0}, O_{t1}, \dots)

Log data, O_d
 Some number of bytes data

Transaction Receipts Trie

Index keyed trie




Ethereum Block Fields

- Header
 - Parent Hash
 - Ommers Hash
 - Beneficiary
 - State Root
 - Transaction Root
 - Receipts Root
 - Logs Bloom
 - Difficulty
 - Block Number
 - Gas Limit
 - Gas Used
 - Timestamp
 - Extra Data
 - Mix Hash
 - Nonce
- Ommer Block Headers
- Transactions

See [Beige Paper](#)

Ommers (Uncle) Blocks

Unlike Bitcoin, Ethereum gives rewards to blocks that have been produced and are valid, but do not form part of the canonical chain. Has to be within the last six blocks



Orphan in Bitcoin (No reward)
Ommers in Ethereum (Small reward)


Ommmer block defined as “The child of an ancestor that is not itself an ancestor”

- Smaller reward given for ommer blocks
- Transactions from ommer blocks are not included

2 Types of accounts in Ethereum

Externally Owned Accounts (EOA)

Controlled by people + private keys



Contract accounts (Smart Contracts)

Controlled by smart contract code + storage



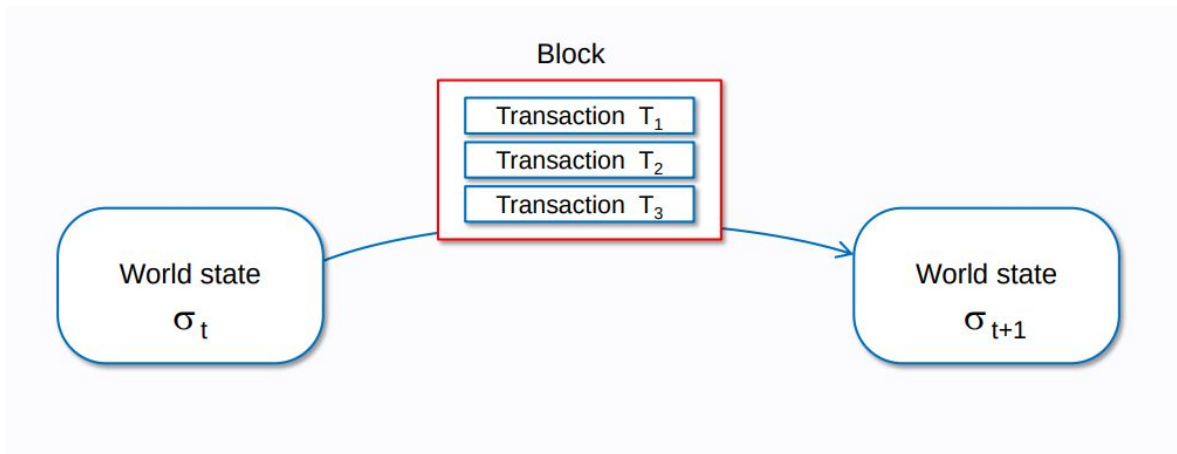
Note that because a contract account does not have a private key, it cannot initiate a transaction. Only EOAs can initiate transactions, but contracts can react to transactions by calling other contracts, building complex execution paths.

The Account State

The account state comprises the following four fields:

- **balance:** A scalar value equal to the number of Wei owned by this address.
- **storageRoot:** A 256-bit hash of the root node of a Merkle Patricia tree that encodes the storage contents of the account (a mapping between 256-bit integer values)
- **codeHash:** The hash of the EVM code of this account
- **nonce:** A scalar value equal to the number of transactions sent from this address or, in the case of accounts with associated code, the number of contract-creations made by this account.

Ethereum Transactions



A submitted transaction includes the following information:

- `recipient` – the receiving address (if an externally-owned account, the transaction will transfer value. If a contract account, the transaction will execute the contract code)
- `signature` – the identifier of the sender. This is generated when the sender's private key signs the transaction and confirms the sender has authorised this transaction
- `value` – amount of ETH to transfer from sender to recipient (in WEI, a denomination of ETH)
- `data` – optional field to include arbitrary data
- `gasLimit` – the maximum amount of gas units that can be consumed by the transaction. Units of gas represent computational steps
- `maxPriorityFeePerGas` - the maximum amount of gas to be included as a tip to the miner
- `maxFeePerGas` - the maximum amount of gas willing to be paid for the transaction (inclusive of `baseFeePerGas` and `maxPriorityFeePerGas`)


Some practical points about transactions

- Miners choose which transactions to include in a block
- Miners can add their own transactions to a block
- Miners choose the order of transactions in a block
- Your transaction is in competition with other transactions for inclusion in the block

Processing a transaction

Initial account balance = 5

line 1
line 2
line 3
line 4
line 5



Process flow

Final account balance = 10

Transactions are atomic

Initial account balance = 5

EVM reverts transaction

line 1
line 2
line 3
line 4
line 5

Final account balance = 5

Process flow



Block Explorer

The Ethereum Blockchain Explorer

[All Filters](#)

ETHER PRICE

\$3,345.60 @ 0.06645 BTC (+2.72%)



TRANSACTIONS

1,256.40 M (15.7 TPS)

MED GAS PRICE

39 Gwei (\$2.74)

ETHEREUM TRANSACTION HISTORY IN 14 DAYS



MARKET CAP

\$392,180,105,051.00



DIFFICULTY

8,075.63 TH

HASH RATE

616,077.75 GH/s



Latest Blocks

Bk	13081315 27 secs ago	Miner firepool 257 txns in 20 secs	3,71611 Eth
Bk	13081313 52 secs ago	Miner BeePool 65 txns in 26 secs	2,03172 Eth
Bk	13081312 1 min 18 secs ago	Miner 0x0708f87a089a91c65d...	2,17696 Eth
Bk	13081311 1 min 50 secs ago	Miner Ethermine 383 txns in 8 secs	2,2938 Eth
Bk	13081310 1 min 58 secs ago	Miner Nanopool 251 txns in 2 secs	2,2438 Eth
Bk	13081309 2 min 0 secs ago	Miner BeePool 82 txns in 10 secs	2,03181 Eth

[View all blocks](#)

Latest Transactions

Tx	0x9f6520ded15e...	From 0x8a689700adbfc2dab7... To 0x198fa2c6f9591ac9fd9...	0.00556 Eth
Tx	0xcb79e81bc04bf...	From 0xcdfa1be78735b56ad9... To 0x2f151af3999d339f3a9...	0.0006 Eth
Tx	0x3cf1e4b736e...	From 0x5d9d4a102bcf2d8ed8... To 0x7a250d5630b4cf5397...	0 Eth
Tx	0x297d2f1d02cf0...	From 0x92573646c0c351f666... To 0x1a2a1c938ce3ec39b6...	0 Eth
Tx	0xa275cfdea754...	From 0xffbcd95ef09e297dc33d... To 0x31c8eacbfdd875c74b...	0 Eth
Tx	0xa4534caf765a...	From 0x2d5427f30e0085d25c... To 0x7a7a75d4f6a30b4e4f5307...	0.2 Eth

[View all transactions](#)

Using the blockchain explorer, have a look at the following transactions, what do they do ?

1. 0xb9316bbbae1cd21cd824de8651c72582261230724b1957abdeb466aa96a359c9
2. 0x4fc1580e7f66c58b7c26881cce0aab9c3509afe6e507527f30566fbf8039bcd0
3. 0x552bc0322d78c5648c5efa21d2daa2d0f14901ad4b15531f1ab5bbe5674de34f
4. 0x7a026bf79b36580bf7ef174711a3de823ff3c93c65304c3acc0323c77d62d0ed
5. 0x814e6a21c8eb34b62a05c1d0b14ee932873c62ef3c8575dc49bcf12004714eda

What is the largest account balance you can find ?

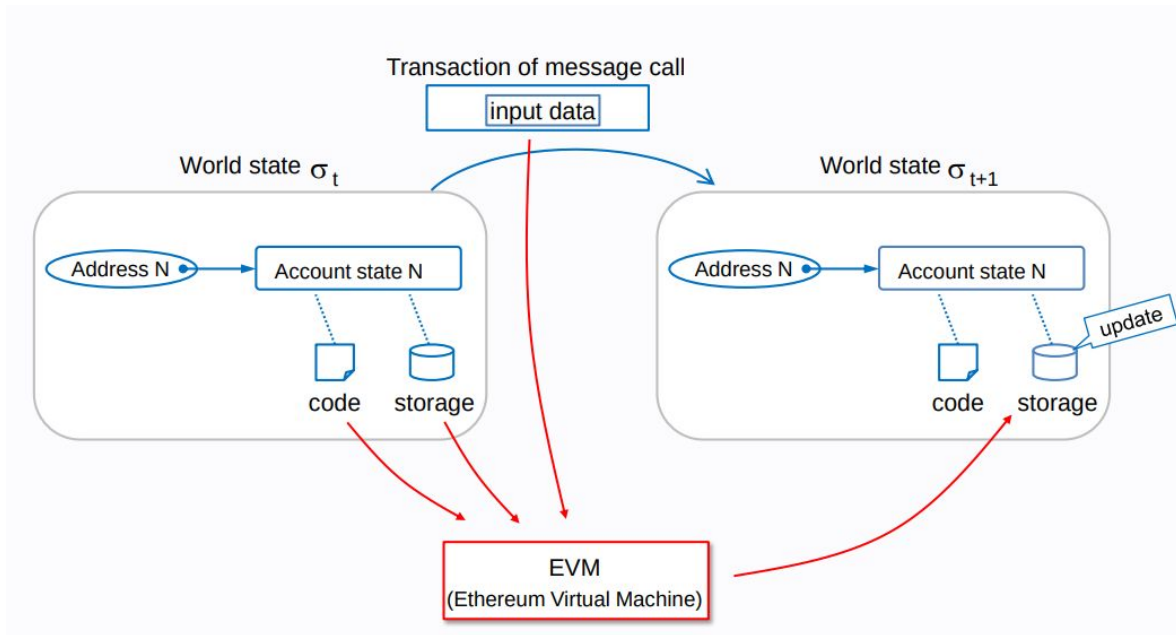
What is special about these accounts :

- 0x1db3439a222c519ab44bb1144fc28167b4fa6ee6
- 0x00000000000000000000000000000000dEaD

Ethereum Virtual Machine

Ethereum is a distributed state machine.

Ethereum's state is a large data structure which holds not only all accounts and balances, but a *machine state*, which can change from block to block according to a pre-defined set of rules, and which can execute arbitrary machine code.



From : [Ethereum EVM](#)

The EVM is a stack machine , the stack has a maximum size of 1024.

Stack items have a size of 256 bits; in fact, the EVM is a 256-bit word machine (this facilitates Keccak256 hash scheme and elliptic-curve computations).

During execution 2 areas are available for variables

- memory - a transient memory which does not persist between transactions
- storage - part of a Merkle Patricia storage trie associated with the contract's account, part of the global state

Execution / OP codes

A compiled smart contract bytecode executes as a number of EVM opcodes

There is a maximum of 256 possible op codes : [Reference sheet](#) for op codes

Op codes can be categorised as

- **Stack-manipulating opcodes** (*POP, PUSH, DUP, SWAP*)
- **Arithmetic/comparison/bitwise opcodes** (*ADD, SUB, GT, LT, AND, OR*)
- **Environmental opcodes** (*CALLER, CALLVALUE, NUMBER*)
- **Memory-manipulating opcodes** (*MLOAD, MSTORE, MSTORE8, MSIZE*)
- **Storage-manipulating opcodes** (*SLOAD, SSTORE*)
- **Program counter related opcodes** (*JUMP, JUMPI, PC, JUMPDEST*)
- **Halting opcodes** (*STOP, RETURN, REVERT, INVALID, SELFDESTRUCT*)

EVM Languages

- Solidity
 - The most popular programming language for Ethereum contracts
- LLL
 - *Low-level Lisp-like Language*
- Vyper
 - A language with overflow-checking, numeric units but without unlimited loops
- Yul / Yul+
 - An intermediate language that can be compiled to bytecode for different backends. Support for EVM 1.0, EVM 1.5 and Ewasm is planned, and it is designed to be a usable common denominator of all three platforms.
- FE
 - Statically typed language Inspired by Rust and Python
- Pyramid Scheme (experimental)
 - A Scheme compiler into EVM that follows the SICP compilation approach
- Flint
 - A language with several security features: e.g. asset types with a restricted set of atomic operations
- LLLL
 - An LLL-like compiler being implemented in Isabelle/HOL
- HAssembly-evm
 - An EVM assembly implemented as a Haskell DSL
- Bamboo (experimental)
 - A language without loops

Ethereum Clients

Name	Language	Sync Strategy	State Pruning
Geth	Go	Fast, Full	Archive, Pruned
Open Ethereum	Rust	Warp, Full	Archive, Pruned
Nethermind	C#	Fast, Full	Archive, Pruned
Besu	Java	Fast, Full	Archive, Pruned
Erigon	Go	Fast, Full	Archive, Pruned

Sync Modes

- Full – downloads all blocks (including headers, transactions and receipts) and generates the state of the blockchain incrementally by executing every block.
- Fast (Default) – downloads all blocks (including headers, transactions and receipts), verifies all headers, and downloads the state and verifies it against the headers.
- Light – downloads all block headers, block data, and verifies some randomly.
- Warp sync – Every 5,000 blocks, nodes will take a consensus-critical snapshot of that block's state. Any node can fetch these snapshots over the network, enabling a fast sync.
- Beam sync – A sync mode that allows you to get going faster. It doesn't require long waits to sync, instead it back-fills data over time.
- Header sync – you can use a trusted checkpoint to start syncing from a more recent header and then leave it up to a background process to fill the gaps eventually

Mining on Ethereum

Miners run Ethash (PoW algorithm)

- Uses a 4 GB dataset, which is updated every 30,000 blocks

Ethash is intended to satisfy the following goals:

1. IO saturation: The algorithm should consume nearly the entire available memory access bandwidth (this is a strategy toward achieving ASIC resistance)
2. GPU friendliness: We try to make it as easy as possible to mine with GPUs. Targeting CPUs is almost certainly impossible, as potential specialization gains are too great, and there do exist criticisms of CPU-friendly algorithms that they are vulnerable to botnets, so we target GPUs as a compromise.
3. Light client verifiability: a light client should be able to verify a round of mining in under 0.01 seconds on a desktop in C, and under 0.1 seconds in Python or Javascript, with at most 1 MB of memory (but exponentially increasing)
4. Light client slowdown: the process of running the algorithm with a light client should be much slower than the process with a full client, to the point that the light client algorithm is not an economically viable route toward making a mining implementation, including via specialized hardware.
5. Light client fast startup: a light client should be able to become fully operational and able to verify blocks within 40 seconds in Javascript.

From [Ethereum Wiki](#)

Ethereum Test Networks

The Ropsten test network is a Proof-of-Work testnet for Ethereum. To acquire ETH on Ropsten, one can mine on the network.

The Kovan test network is a Proof-of-Authority testnet for Ethereum, originally started by the Parity team. To acquire ETH on Kovan, one can request it from a faucet.

The Rinkeby test network is a Proof-of-Authority testnet for Ethereum, originally started by the Geth team. To acquire ETH on Rinkeby, one can request it from a faucet.

The Görli test network is a Proof-of-Authority testnet for Ethereum, originally proposed by Chainsafe and Afri Schoedon. To acquire ETH on Görli, one can use the one-way throttled bridge from any of the other three test networks.

Exercise 2

Get some test ether from the Rinkeby Faucet

Send 0.13 ETH to a colleague (Double check that it is the on Rinkeby)

Can you find your transaction on a block explorer ?

Was the gas estimate accurate ?

Did you get the expected gas price ?

Next lesson

Introduction to Remix